

Homework #0

Spring 2020, CSE 446/546: Machine Learning

Dino Bektsevic

Probability and Statistics

A.1 [2 points] (Bayes Rule, from Murphy exercise 2.4.) After your yearly checkup, the doctor has bad news and good news. The bad news is that you tested positive for a serious disease, and that the test is 99% accurate (i.e., the probability of testing positive given that you have the disease is 0.99, as is the probability of testing negative given that you don't have the disease). The good news is that this is a rare disease, striking only one in 10,000 people. What are the chances that you actually have the disease? (Show your calculations as well as giving the final result.)

Denoting $X = 1$ as test result being positive and $Y = 1$ as actually having the disease, following Bayes rule we have

$$\begin{aligned} P(Y = 1|X = 1) &= \frac{P(X = 1|Y = 1)P(Y = 1)}{P(X = 1|Y = 1) + P(X = 1|Y = 0)P(Y = 0)} \\ &= \frac{0.99 \cdot 0.0001}{0.99 \cdot 0.0001 + 0.01 \cdot 0.9999} \\ &= 0.0098 \approx 1\% \end{aligned}$$

since we are given that $P(X = 1|Y = 1) = 0.99$ and $P(Y = 1) = 1/10000 = 0.0001$ from which it follows that $P(Y = 0) = 1 - P(Y = 1) = 0.9999$ and $P(X = 1|Y = 0) = 1 - P(X = 1|Y = 1) = 0.01$

A.2 For any two random variables X, Y the *covariance* is defined as $\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$. You may assume X and Y take on a discrete values if you find that is easier to work with.

- a. [1 points] If $\mathbb{E}[Y|X = x] = x$ show that $\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])^2]$. Simplify the right hand side of the expression for covariance (marking expectation values that are scalars by replacing them with μ to show they can be taken out of expectation value operator):

$$\begin{aligned} \mathbb{E}[(X - \mathbb{E}[X])^2] &= \mathbb{E}[X^2 - 2XE(X) + E(X)^2] = \mathbb{E}[X^2 - 2\mu_x X + 2\mu_x^2] \\ &= \mathbb{E}[X^2] - 2\mu_x \mathbb{E}[X] + \mu_x^2 = \mathbb{E}[X^2] - 2\mu_x^2 + \mu_x^2 \\ &= \mathbb{E}[X^2] - \mathbb{E}[X]^2 = \text{Var}(X) \end{aligned}$$

It's clear now that we are actually looking to show that $\text{Cov}(X, Y) = \text{Var}(X)$. Apply what is given in the problem to law of iterative expectations:

$$\mathbb{E}[Y] = \int yp(y)dy = \int \mathbb{E}(Y|X = x)p(x)dx = \int xp(x)dx = \mathbb{E}[X]$$

Finally, write out the expression for covariance of two rvs:

$$\text{Cov}(X, Y) = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y] = \mathbb{E}[XY] - \mathbb{E}[X]^2 = \text{Var}(X)$$

since we can write the joint probability $\mathbb{E}[XY] = \int \int xp(x)yp_{Y|X=x}(y)dxdy = \int xp(x)xdx = \int x^2p(x)dx = \mathbb{E}[X^2]$

- b. [1 points] If X, Y are independent show that $\text{Cov}(X, Y) = 0$.

$$\text{Cov}(X, Y) = \mathbb{E}(X, Y) - \mathbb{E}(X)\mathbb{E}(Y) = \mathbb{E}(X)\mathbb{E}(Y) - \mathbb{E}(X)\mathbb{E}(Y) = 0$$

A.3 Let X and Y be independent random variables with PDFs given by f and g , respectively. Let h be the PDF of the random variable $Z = X + Y$.

- a. [2 points] Show that $h(z) = \int_{-\infty}^{\infty} f(x)g(z-x)dx$. (If you are more comfortable with discrete probabilities, you can instead derive an analogous expression for the discrete case, and then you should give a one sentence explanation as to why your expression is analogous to the continuous case.).
- b. [1 points] If X and Y are both independent and uniformly distributed on $[0, 1]$ (i.e. $f(x) = g(x) = 1$ for $x \in [0, 1]$ and 0 otherwise) what is h , the PDF of $Z = X + Y$?

A.4 [1 points] A random variable $X \sim \mathcal{N}(\mu, \sigma^2)$ is Gaussian distributed with mean μ and variance σ^2 . Given that for any $a, b \in \mathbb{R}$, we have that $Y = aX + b$ is also Gaussian, find a, b such that $Y \sim \mathcal{N}(0, 1)$.

$$\begin{aligned} E[Y] &= aE[X] + b = 0 \\ \text{Var}[Y] &= a^2\text{Var}[X] = 1 \\ a\mu + b &= 0 \\ a^2\sigma^2 &= 1 \\ a &= 1/\sigma \\ b &= -\mu/\sigma \end{aligned}$$

A.5 [2 points] For a random variable Z , its mean and variance are defined as $\mathbb{E}[Z]$ and $\mathbb{E}[(Z - \mathbb{E}[Z])^2]$, respectively. Let X_1, \dots, X_n be independent and identically distributed random variables, each with mean μ and variance σ^2 . If we define $\hat{\mu}_n = \frac{1}{n} \sum_{i=1}^n X_i$, what is the mean and variance of $\sqrt{n}(\hat{\mu}_n - \mu)$?

The average value of many random i.i.d. variables that have the same mean μ is the mean μ itself. So :

$$\hat{\mu}_n = \frac{1}{n} \sum X_i = \mu$$

so:

$$\sqrt{n}(\hat{\mu}_n - \mu) = 0$$

The variance of many i.i.d. rvs can be written as:

$$\begin{aligned} \text{Var}[\sqrt{n}\hat{\mu}_n - \sqrt{n}\mu] &= n\text{Var}[\hat{\mu}_n] + n\text{Var}[\mu] \\ &= n\text{Var}[\hat{\mu}_n] = \frac{n\sigma^2}{n} = \sigma^2 \end{aligned}$$

since $\text{Var}[aX] = a^2\text{Var}[X]$ and $\text{Var}[\mu] = 0$ because all observed rvs have the same mean.

A.6 If $f(x)$ is a PDF, the cumulative distribution function (CDF) is defined as $F(x) = \int_{-\infty}^x f(y)dy$. For any function $g : \mathbb{R} \rightarrow \mathbb{R}$ and random variable X with PDF $f(x)$, recall that the expected value of $g(X)$ is defined as $\mathbb{E}[g(X)] = \int_{-\infty}^{\infty} g(y)f(y)dy$. For a boolean event A , define $\mathbf{1}\{A\}$ as 1 if A is true, and 0 otherwise. Thus, $\mathbf{1}\{x \leq a\}$ is 1 whenever $x \leq a$ and 0 whenever $x > a$. Note that $F(x) = \mathbb{E}[\mathbf{1}\{X \leq x\}]$. Let X_1, \dots, X_n be independent and identically distributed random variables with CDF $F(x)$. Define $\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}\{X_i \leq x\}$. Note, for every x , that $\hat{F}_n(x)$ is an empirical estimate of $F(x)$. You may use your answers to the previous problem.

- a. [1 points] For any x , what is $\mathbb{E}[\hat{F}_n(x)]$?
- b. [1 points] For any x , the variance of $\hat{F}_n(x)$ is $\mathbb{E}[(\hat{F}_n(x) - F(x))^2]$. Show that $\text{Variance}(\hat{F}_n(x)) = \frac{F(x)(1-F(x))}{n}$.

- c. [1 points] Using your answer to b, show that for all $x \in \mathbb{R}$, we have $\mathbb{E}[(\hat{F}_n(x) - F(x))^2] \leq \frac{1}{4n}$.

B.1 [1 points] Let X_1, \dots, X_n be n independent and identically distributed random variables drawn uniformly at random from $[0, 1]$. If $Y = \max\{X_1, \dots, X_n\}$ then find $\mathbb{E}[Y]$.

Linear Algebra and Vector Calculus

A.7 (Rank) Let $A = \begin{bmatrix} 1 & 2 & 1 \\ 1 & 0 & 3 \\ 1 & 1 & 2 \end{bmatrix}$ and $B = \begin{bmatrix} 1 & 2 & 3 \\ 1 & 0 & 1 \\ 1 & 1 & 2 \end{bmatrix}$. For each matrix A and B ,

- a. [2 points] what is its rank?

Rank is just the number of linearly independent columns of the matrix:

$$\text{rank}(A) = 3$$

$$\text{rank}(B) = 2$$

For the B matrix if we take away the last row from the first we are left with $[0, 1, 1]$ and if we take away second row from the last we are left again with $[0, 1, 1]$ so at least one of the rows is not linearly independent from the remaining two.

- b. [2 points] what is a (minimal size) basis for its column span?

All columns of matrix A are linearly independent, so the basis of the image space is subset of \mathbb{R}^3 stretched by the set of column vectors:

$$\left\{ \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \begin{bmatrix} 2 \\ 0 \\ 1 \end{bmatrix}, \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix} \right\}$$

Only two columns of matrix B are linearly independent. Starting from the note in previous problem:

$$A \approx \begin{bmatrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

from which we see that first column can be written as $c_1 = c_3 - c_2$. So the basis of the image space is a subset of \mathbb{R}^2 stretched by column vectors:

$$\left\{ \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} \right\}$$

A.8 (Linear equations) Let $A = \begin{bmatrix} 0 & 2 & 4 \\ 2 & 4 & 2 \\ 3 & 3 & 1 \end{bmatrix}$, $b = [-2 \quad -2 \quad -4]^T$, and $c = [1 \quad 1 \quad 1]^T$.

- a. [1 points] What is Ac ?

It's a matrix multiplication operation!

$$\begin{bmatrix} 0 & 2 & 4 \\ 2 & 4 & 2 \\ 3 & 3 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 2+4 \\ 2+4+2 \\ 3+3+1 \end{bmatrix} = \begin{bmatrix} 6 \\ 8 \\ 7 \end{bmatrix}$$

- b. [2 points] What is the solution to the linear system $Ax = b$? (Show your work).

Sequence of operations is as follows:

- (a) switch first and second row and divide first row by 2 (sets (1,1) to 1) and zero out first column,
- (b) divide second row by two and zero out second column,
- (c) divide 3rd row by 4 and zero out 3rd column.

$$\begin{bmatrix} 0 & 2 & 4 \\ 2 & 4 & 2 \\ 3 & 3 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} -2 \\ -2 \\ -4 \end{bmatrix}$$


$$\left[\begin{array}{ccc|c} 0 & 2 & 4 & -2 \\ 2 & 4 & 2 & -2 \\ 3 & 3 & 1 & -4 \end{array} \right] \approx \left[\begin{array}{ccc|c} 1 & 2 & 1 & -2 \\ 0 & 2 & 4 & -2 \\ 0 & -3 & -2 & -1 \end{array} \right] \approx$$

$$\left[\begin{array}{ccc|c} 1 & 0 & 1 & -3 \\ 0 & 1 & 2 & -1 \\ 0 & 0 & 4 & -4 \end{array} \right] \approx \left[\begin{array}{ccc|c} 1 & 0 & 0 & -2 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & -1 \end{array} \right] \rightarrow x = \begin{bmatrix} -2 \\ 1 \\ -1 \end{bmatrix}$$

A.9 (Hyperplanes) Assume w is an n -dimensional vector and b is a scalar. A hyperplane in \mathbb{R}^n is the set $\{x : x \in \mathbb{R}^n, \text{ s.t. } w^T x + b = 0\}$.

- a. [1 points] ($n = 2$ example) Draw the hyperplane for $w = [-1, 2]^T$, $b = 2$? Label your axes. Effectively this gives us an equation of a line $-x_1 + 2x_2 + 2 = 0$ or $x_1 = 2x_2 + 2$. Plot line using following Python code:

```
import matplotlib.pyplot as plt
import numpy as np
x2 = np.arange(-10, 10, 1)
x1 = 2*x2+2
plt.plot(x1, x2)
plt.xlabel("x2")
plt.ylabel("x1")
plt.grid()
```




- b. [1 points] ($n = 3$ example) Draw the hyperplane for $w = [1, 1, 1]^T$, $b = 0$? Label your axes. Following the same principles above $x + y + z = 0 \rightarrow z = -x - y$ we have:

```
import matplotlib.pyplot as plt
from mpl_toolkits import mplot3d
import numpy as np

x2 = np.arange(-10, 10, 1)
x3 = np.arange(-10, 10, 1)
X2, X3 = np.meshgrid(x2, x3)
X1 = -X2 -X3

fig = plt.figure()
ax = plt.axes(projection='3d')
ax.contour3D(X1, X2, X3, 150)
ax.set_zlabel("x3")
ax.set_xlabel("x2")
ax.set_ylabel("x1")
plt.grid()
```



- c. [2 points] Given some $x_0 \in \mathbb{R}^n$, find the *squared distance* to the hyperplane defined by $w^T x + b = 0$. In other words, solve the following optimization problem:

$$\begin{aligned} \min_x & \|x_0 - x\|^2 \\ \text{s.t. } & w^T x + b = 0 \end{aligned}$$

(Hint: if \tilde{x}_0 is the minimizer of the above problem, note that $\|x_0 - \tilde{x}_0\| = \left| \frac{w^T(x_0 - \tilde{x}_0)}{\|w\|} \right|$. What is $w^T \tilde{x}_0$?)
We are looking to minimize the following function:

$$\begin{aligned} f(\vec{x}) &= \omega^T \|\vec{x} - \vec{x}_0\|^2 + b = 0 \\ &= \omega^T \|\vec{x}^T \vec{x} - 2\vec{x} \cdot \vec{x}_0 + \vec{x}_0^T \vec{x}_0\| + b \\ &= \omega \vec{x}^T \vec{x} - 2\omega^T \|\vec{x} \cdot \vec{x}_0\| + \omega^T \vec{x}_0^T \vec{x}_0 + b = 0 \end{aligned}$$

Note that values $\vec{x}^T \vec{x}$ and $\vec{x}_0^T \vec{x}_0$ are just scalars. Minimization can be performed via the method of Lagrange multipliers of the above equation given the constraint that \vec{x} lies in the plane $\Phi = \omega^T \vec{x} + b = 0$:

$$\begin{aligned} \frac{\partial f(\vec{x})}{\partial x_i} - \lambda \frac{\partial \Phi(\vec{x})}{\partial x_i} &= 0 \\ -2\omega^T \frac{\partial \|\vec{x} \cdot \vec{x}_0\|}{\partial x_i} - \lambda \omega^T \frac{\partial \vec{x}}{\partial x_i} &= 0 \\ -2\omega^T \frac{\partial \vec{x}}{\partial x_i} \cdot \vec{x}_0 - \lambda \omega^T \frac{\partial \vec{x}}{\partial x_i} &= 0 \end{aligned}$$

Following is speculative

$$(-2\omega^T \vec{x}_0 - \lambda \omega^T) \frac{\partial \vec{x}}{\partial x_i} = 0$$

Note that $\omega^T \vec{x}_0$ is a scalar, that x_0 denotes the magnitude of the radius vector of the point we're measuring distance to so it also is a scalar, we can

A.10 For possibly non-symmetric $\mathbf{A}, \mathbf{B} \in \mathbb{R}^{n \times n}$ and $c \in \mathbb{R}$, let $f(x, y) = x^T \mathbf{A} x + y^T \mathbf{B} x + c$. Define $\nabla_z f(x, y) = \left[\frac{\partial f(x, y)}{\partial z_1} \quad \frac{\partial f(x, y)}{\partial z_2} \quad \dots \quad \frac{\partial f(x, y)}{\partial z_n} \right]^T$.

- [2 points] Explicitly write out the function $f(x, y)$ in terms of the components $A_{i,j}$ and $B_{i,j}$ using appropriate summations over the indices.
- [2 points] What is $\nabla_x f(x, y)$ in terms of the summations over indices *and* vector notation?

c. [2 points] What is $\nabla_y f(x, y)$ in terms of the summations over indices *and* vector notation?

B.2 [1 points] The *trace* of a matrix is the sum of the diagonal entries; $\text{Tr}(A) = \sum_i A_{ii}$. If $A \in \mathbb{R}^{n \times m}$ and $B \in \mathbb{R}^{m \times n}$, show that $\text{Tr}(AB) = \text{Tr}(BA)$.

It can be shown in general that trace is invariant under cyclic permutations, which for case of $n=2$ looks like commutation:

$$\text{tr}(AB) = \sum_i (ab)_{ii} = \sum_{i=1}^m \sum_{j=1}^n a_{ij} b_{ji} = \sum_{j=1}^n \sum_{i=1}^m b_{ji} a_{ij} = \sum_i (ab)_{jj} = \text{tr}(BA)$$

using the fact that product of $n \times m$ and $m \times n$ matrix is an $m \times m$ matrix and that the trace of a square matrix can be rewritten as:

$$\text{tr}(A^T B) = \sum_{i,j} A_{ij} B_{ij}$$

as per Wikipedia.

B.3 [1 points] Let v_1, \dots, v_n be a set of non-zero vectors in \mathbb{R}^d . Let $V = [v_1, \dots, v_n]$ be the vectors concatenated.

a. What is the minimum and maximum rank of $\sum_{i=1}^n v_i v_i^T$?

Basis of an \mathbb{R}^d is spanned by d linearly independent vectors. So the maximum rank of a matrix will be d . Minimal rank of a matrix would always be 0 in case of a zero matrix.

b. What is the minimum and maximum rank of V ?

Same as a) (are these supposed to be like trick questions?)

c. Let $A \in \mathbb{R}^{D \times d}$ for $D > d$. What is the minimum and maximum rank of $\sum_{i=1}^n (Av_i)(Av_i)^T$?

d. What is the minimum and maximum rank of AV ? What if V is rank d ?

Programming

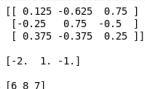
A.11 For the A, b, c as defined in Problem 8, use NumPy to compute (take a screen shot of your answer):

a. [2 points] What is A^{-1} ?

b. [1 points] What is $A^{-1}b$? What is Ac ?

```
In [51]: import numpy as np
A = np.array([[0,2,4], [2,4,2], [3,3,1]])
b = np.array([-2, -2, -4])
c = np.array([1, 1, 1])

print(np.linalg.inv(A))
print()
print(np.linalg.inv(A)@b)
print()
print(A@c)
```



```
[[ 0.125 -0.625  0.75 ]
 [-0.25  0.75 -0.5 ]
 [ 0.375 -0.375  0.25 ]]

[-2.  1. -1.]

[6 8 7]
```

A.12 [4 points] Two random variables X and Y have equal distributions if their CDFs, F_X and F_Y , respectively, are equal, i.e. for all x , $|F_X(x) - F_Y(x)| = 0$. The central limit theorem says that the sum of k independent, zero-mean, variance-1/ k random variables converges to a (standard) Normal distribution as k goes off to infinity. We will study this phenomenon empirically (you will use the Python packages Numpy and Matplotlib). Define $Y^{(k)} = \frac{1}{\sqrt{k}} \sum_{i=1}^k B_i$ where each B_i is equal to -1 and 1 with equal probability. From your solution to problem 5, we know that $\frac{1}{\sqrt{k}} B_i$ is zero-mean and has variance $1/k$.

a. For $i = 1, \dots, n$ let $Z_i \sim \mathcal{N}(0, 1)$. If $F(x)$ is the true CDF from which each Z_i is drawn (i.e., Gaussian) and $\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}\{Z_i \leq x\}$, use the answer to problem 1.5 above to choose n large enough such that, for all $x \in \mathbb{R}$, $\sqrt{\mathbb{E}[(\hat{F}_n(x) - F(x))^2]} \leq 0.0025$, and plot $\hat{F}_n(x)$ from -3 to 3 . (Hint: use `Z=numpy.random.randn(n)` to generate the random variables, and `import matplotlib.pyplot as plt; plt.step(sorted(Z), np.arange(1,n+1)/float(n))` to plot).


```

import matplotlib.pyplot as plt
import numpy as np

n = 20000
Z = np.sign(np.random.randn(n))
plt.step(sorted(Z), np.arange(1, n + 1) / float(n))

plt.xlim(-3, 3)
plt.xlabel("Observations")
plt.ylabel("Probability")
plt.show()

```



- b. For each $k \in \{1, 8, 64, 512\}$ generate n independent copies $Y^{(k)}$ and plot their empirical CDF on the same plot as part a.
(Hint: `np.sum(np.sign(np.random.randn(n, k))*np.sqrt(1./k), axis=1)` generates n of the $Y^{(k)}$ random variables.)

```

import matplotlib.pyplot as plt
import numpy as np

def Yk(n, k=1):
    """Returns and array of samples of function:

        Y^{(k)} = 1/sqrt(k) \sum_{i=1}^k B_i

    where B_i = +1 or -1 with equal probability.

    Parameters
    -----
    n : 'int'
        Number of samples of Y^{(k)}.
    k : 'int'
        Upper bound of the Y^{(k)} sum
    """
    B = np.sign(np.random.randn(n, k))
    return np.sum(np.sqrt(1.0/k)*B, axis=1)

n = 20000
for k in [1, 8, 64, 512]:
    Y = Yk(n, k)
    plt.step(sorted(Y), np.arange(1, n + 1) / float(n), label='{0}'.format(k))

G = np.random.normal(size=n)
plt.step(sorted(G), np.arange(1, n + 1) / float(n), label='Gaussian')

plt.legend()
plt.xlim(-3, 3)
plt.xlabel("Observations")
plt.ylabel("Probability")
plt.show()

```

